

In [1]:

```
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
import numpy as np
import pandas as pd
```

In [3]:

```
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

In [4]:

```
car=pd.read_csv('D:/task/carprice.csv')
car.head()
```

Out[4]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	eng
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns

In [5]:

```
car.shape
```

Out[5]:

(205, 26)

In [6]:

car.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel            205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength              205 non-null    float64
11  carwidth               205 non-null    float64
12  carheight              205 non-null    float64
13  curbweight             205 non-null    int64
14  enginetype             205 non-null    object
15  cylindernumber         205 non-null    object
16  enginesize             205 non-null    int64
17  fuelsystem             205 non-null    object
18  boreratio              205 non-null    float64
19  stroke                 205 non-null    float64
20  compressionratio       205 non-null    float64
21  horsepower             205 non-null    int64
22  peakrpm                205 non-null    int64
23  citympg                205 non-null    int64
24  highwaympg             205 non-null    int64
25  price                  205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

In [7]:

car.describe()

Out[7]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	e
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	2
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	1
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	1
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	1
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	3

In [8]:

```
cars=cars.select_dtypes(include=['float64','int64'])
cars.head()
```

Out[8]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	borer
0	1	3	88.6	168.8	64.1	48.8	2548	130	:
1	2	3	88.6	168.8	64.1	48.8	2548	130	:
2	3	1	94.5	171.2	65.5	52.4	2823	152	:
3	4	2	99.8	176.6	66.2	54.3	2337	109	:
4	5	2	99.4	176.6	66.4	54.3	2824	136	:

In [9]:

```
# dropping symboling and car_ID as symboling is more of categorical variable as described b
#an index type variable and not a predictor
cars= cars.drop(['symboling', 'car_ID'], axis=1)
cars.head()
```

Out[9]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine location	wheelb
0	alfa-romero giulia	gas	std	two	convertible	rwd	front	{
1	alfa-romero stelvio	gas	std	two	convertible	rwd	front	{
2	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	{
3	audi 100 ls	gas	std	four	sedan	fwd	front	{
4	audi 100ls	gas	std	four	sedan	4wd	front	{

5 rows × 24 columns

In [10]:

```
car['symboling'].astype('category').value_counts()
```

Out[10]:

```
0    67
1    54
2    32
3    27
-1   22
-2    3
Name: symboling, dtype: int64
```

In [11]:


```
car['aspiration'].astype('category').value_counts()
```

Out[11]:

```
std      168
turbo     37
Name: aspiration, dtype: int64
```


In [12]:

```
sns.distplot(car['carheight'])
plt.show()
```



In [13]:

```
sns.distplot(car['carwidth'])
plt.show()
```

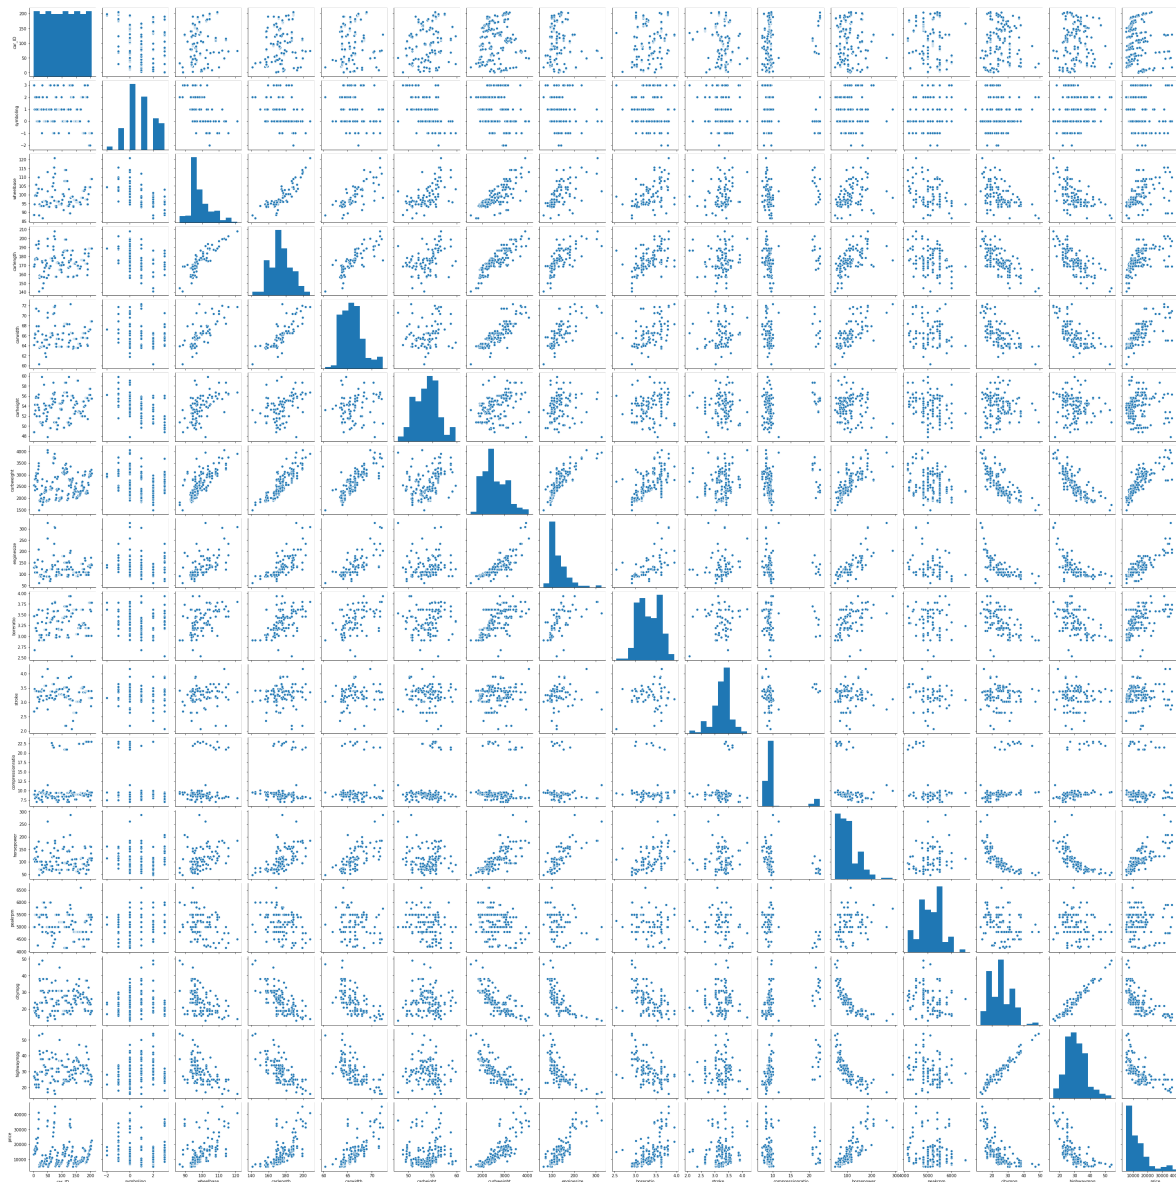


In [14]:

```
#VISUALISING THE DATA
```


```
sns.pairplot(car)
```

```
plt.show()
```



In [15]:

```
for i, col in enumerate(car.columns):
    plt.figure(i)
    sns.scatterplot(x=car[col],y=car['stroke'])
```

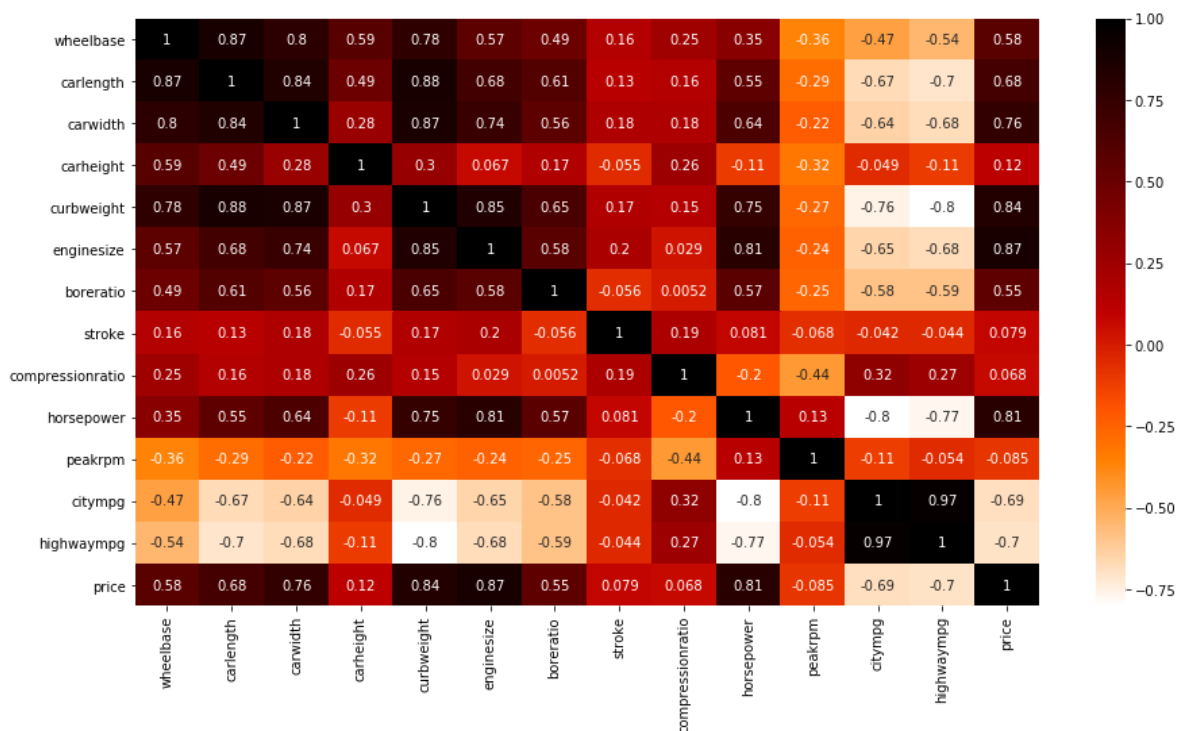


In [16]:

```
#corealtion with Dependent var and independent var's
corr=cars.corr()
plt.figure(figsize=(15,8))
sns.heatmap(corr,annot=True,cmap="gist_heat_r")
```

Out[16]:

<matplotlib.axes._subplots.AxesSubplot at 0x1ede939e1c0>



In [17]:

```
##ONLY CAR NAMES
```

```
carnames = car['CarName'].apply(lambda x: x.split(" ")[0])  
carnames[:21]
```

Out[17]:

```
0    alfa-romero  
1    alfa-romero  
2    alfa-romero  
3         audi  
4         audi  
5         audi  
6         audi  
7         audi  
8         audi  
9         audi  
10        bmw  
11        bmw  
12        bmw  
13        bmw  
14        bmw  
15        bmw  
16        bmw  
17        bmw  
18    chevrolet  
19    chevrolet  
20    chevrolet
```

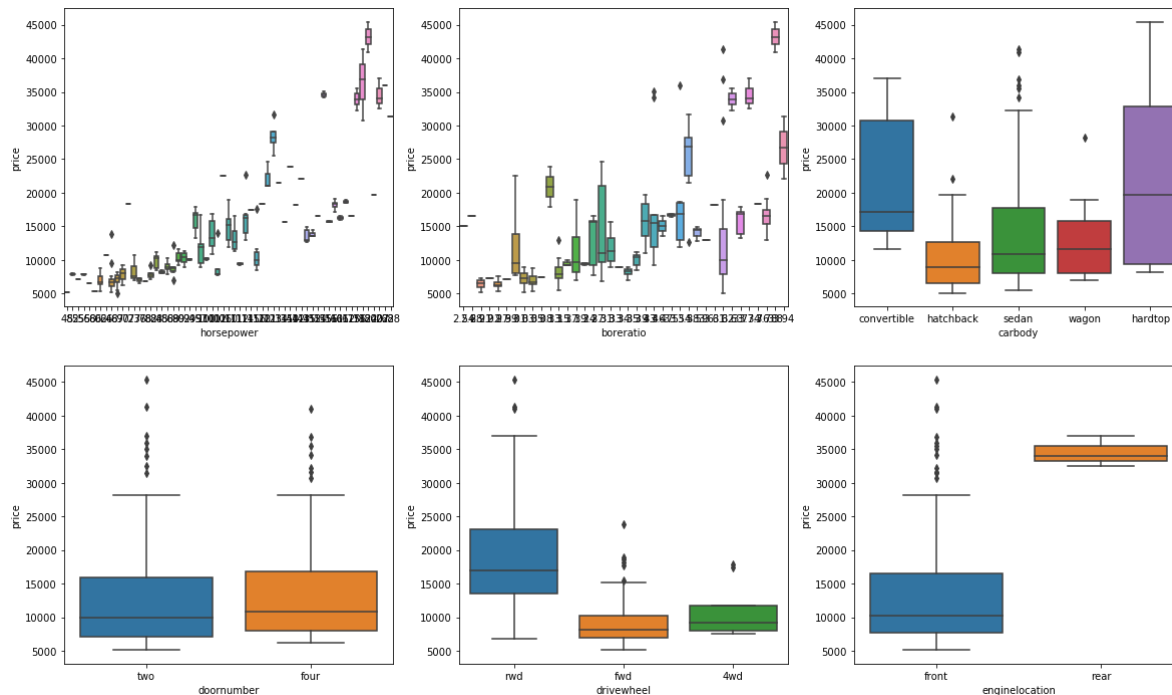
Name: CarName, dtype: object

In [18]:

```


plt.figure(figsize=(20, 12))
plt.subplot(2,3,1)
sns.boxplot(x = 'horsepower', y = 'price', data = car)
plt.subplot(2,3,2)
sns.boxplot(x = 'boreratio', y = 'price', data = car)
plt.subplot(2,3,3)
sns.boxplot(x = 'carbody', y = 'price', data = car)
plt.subplot(2,3,4)
sns.boxplot(x = 'doornumber', y = 'price', data = car)
plt.subplot(2,3,5)
sns.boxplot(x = 'drivewheel', y = 'price', data = car)
plt.subplot(2,3,6)
sns.boxplot(x = 'enginelocation', y = 'price', data = car)
plt.show()

```



In [19]:

```
plt.figure(figsize = (10, 5))  
sns.boxplot(x = 'fueltype', y = 'price', hue = 'drivewheel', data = car)  
plt.show()
```



In [20]:

```
car['symboling'] = car['symboling'].astype('object')
car.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling             205 non-null   object
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(7), object(11)
memory usage: 41.8+ KB
```

In [21]:

```
car['car_company']=carnames  
car['car_company'].value_counts()
```

Out[21]:

toyota	31
nissan	17
mazda	15
honda	13
mitsubishi	13
subaru	12
volvo	11
peugeot	11
dodge	9
volkswagen	9
bmw	8
buick	8
audi	7
plymouth	7
saab	6
isuzu	4
porsche	4
jaguar	3
alfa-romero	3
chevrolet	3
renault	2
vw	2
maxda	2
porcshce	1
vokswagen	1
Nissan	1
toyouta	1
mercury	1

Name: car_company, dtype: int64

In [22]:

```
#bmw
car.loc[(car['car_company']=="bmw"), "car_company"]="BMW"

#toyota
car.loc[(car['car_company']=="toyouta"), "car_company"]="toyota"

# nissan
car.loc[car['car_company'] == "Nissan", 'car_company'] = 'nissan'

# mazda
car.loc[car['car_company'] == "audi", 'car_company'] = 'Audi'

car['car_company'].value_counts()
```

Out[22]:

toyota	32
nissan	18
mazda	15
mitsubishi	13
honda	13
subaru	12
volvo	11
peugeot	11
volkswagen	9
dodge	9
BMW	8
buick	8
plymouth	7
Audi	7
saab	6
isuzu	4
porsche	4
jaguar	3
alfa-romero	3
chevrolet	3
vw	2
renault	2
maxda	2
porcshce	1
vokswagen	1
mercury	1

Name: car_company, dtype: int64

In [23]:

#DATA PREPARATION

```
x=car.drop(columns=['price', "car_ID"])
y=car['price']
y.head()
```

Out[23]:

```
0    13495.0
1    16500.0
2    16500.0
3    13950.0
4    17450.0
Name: price, dtype: float64
```

In [24]:

```
cars_category = x.select_dtypes(include=['object'])
cars_category.head()
```

Out[24]:

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocal
0	3	alfa-romero giulia	gas	std	two	convertible	rwd	f
1	3	alfa-romero stelvio	gas	std	two	convertible	rwd	f
2	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	f
3	2	audi 100 ls	gas	std	four	sedan	fwd	f
4	2	audi 100ls	gas	std	four	sedan	4wd	f

In [25]:

```
cars_dummy = pd.get_dummies(cars_category, drop_first=True)
cars_dummy.head()
```

Out[25]:

	symboling_-1	symboling_0	symboling_1	symboling_2	symboling_3	CarName_alfa-romero Quadrifoglio	CarNan romer
0	0	0	0	0	1	0	
1	0	0	0	0	1	0	
2	0	0	1	0	0	1	
3	0	0	0	1	0	0	
4	0	0	0	1	0	0	

5 rows × 205 columns

In [26]:

```
x=x.drop(columns=cars_category)
x.head()
```

Out[26]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compr
0	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	
1	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	
2	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	
3	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	
4	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	

In [27]:

```
x=pd.concat([x,cars_dummy],axis=1)
x.head()
```

Out[27]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compr
0	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	
1	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	
2	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	
3	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	
4	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	

5 rows × 218 columns

In [28]:

```
x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Columns: 218 entries, wheelbase to car_company_vw
dtypes: float64(7), int64(6), uint8(205)
memory usage: 62.0 KB
```

In [29]:

`x.columns`

Out[29]:

```
Index(['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight',
      'enginesize', 'boreratio', 'stroke', 'compressionratio', 'horsepower',
      ...,
      'car_company_porcshce', 'car_company_porsche', 'car_company_renault',
      'car_company_saab', 'car_company_subaru', 'car_company_toyota',
      'car_company_volkswagen', 'car_company_volkswagen', 'car_company_volv',
      'car_company_vw'],
      dtype='object', length=218)
```

In [30]:

```
#TRAIN-TEST
from sklearn.model_selection import train_test_split

# We specify this so that the train and test data set always have the same rows, respective
np.random.seed(0)
x_train, y_test = train_test_split(cars, train_size = 0.7, test_size = 0.3, random_state =
```

In [31]:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

In [32]:

```
num_vars = ['curbweight', 'carlength', 'curbweight', 'enginesize', 'horsepower', 'price']
x_train[num_vars] = scaler.fit_transform(x_train[num_vars])
```


In [33]:

```
x_train.head()
```

Out[33]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase
122	plymouth fury gran sedan	gas	std	four	sedan	fwd	front	112.0
125	porsche macan	gas	std	two	hatchback	rwd	front	98.5
166	toyota corolla tercel	gas	std	two	hatchback	rwd	front	94.5
1	alfa- romero stelvio	gas	std	two	convertible	rwd	front	101.2
199	volvo diesel	gas	turbo	four	wagon	rwd	front	115.6

5 rows × 24 columns



In [34]:

```
x_train.describe()
```


Out[34]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	
count	143.000000	143.000000	143.000000	143.000000	143.000000	143.000000	143.000000	14
mean	98.523077	0.525476	65.839860	53.551748	0.407878	0.241351	3.307413	
std	5.961835	0.204848	2.214203	2.433766	0.211269	0.154619	0.260997	
min	86.600000	0.000000	60.300000	47.800000	0.000000	0.000000	2.680000	
25%	94.500000	0.399187	63.950000	51.800000	0.245539	0.135849	3.065000	
50%	96.500000	0.502439	65.400000	53.700000	0.355702	0.184906	3.310000	
75%	101.200000	0.669919	66.900000	55.350000	0.559542	0.301887	3.540000	
max	115.600000	1.000000	72.300000	59.100000	1.000000	1.000000	3.940000	




In [35]:

```
plt.figure(figsize = (20, 16))
sns.heatmap(x_train.corr(), annot = True, cmap="YlGnBu")
plt.show()
```



In [36]:

```
plt.figure(figsize=[4,4])
plt.scatter(x_train.carlength, x_train.wheelbase)
plt.show()
```



In [37]:

```
y_train = x_train.pop('stroke')
x_train = x_train
```

In [38]:

```
import statsmodels.api as sm
```

In [39]:

```
#add a constant

x_train_lm = sm.add_constant(x_train[['carlength']])

#create a first fitted model

lr = sm.OLS(y_train,x_train_lm).fit()
```

In [40]:


```
lr.params
```

Out[40]:

```
const      3.108318  
carlength  0.269760  
dtype: float64
```

In [41]:

```
plt.scatter(x_train_lm.iloc[:, 1], y_train)  
plt.plot(x_train_lm.iloc[:, 1], 85.09 + 25.5*x_train_lm.iloc[:, 1], 'r')  
plt.show()
```



In [42]:

```
print(lr.summary())
```

```

=====
                        OLS Regression Results
=====
==
Dep. Variable:          stroke    R-squared:                0.0
31
Model:                  OLS      Adj. R-squared:            0.0
24
Method:                 Least Squares    F-statistic:          4.5
50
Date:                   Sun, 26 Apr 2020    Prob (F-statistic):    0.03
46
Time:                   09:11:23    Log-Likelihood:       -33.8
18
No. Observations:       143    AIC:                  71.
64
Df Residuals:           141    BIC:                  77.
56
Df Model:                1
Covariance Type:        nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.975
const	3.1083	0.071	43.601	0.000	2.967	3.250
carlength	0.2698	0.126	2.133	0.035	0.020	0.519

```

=====
==
Omnibus:                16.116    Durbin-Watson:          2.0
81
Prob(Omnibus):           0.000    Jarque-Bera (JB):       22.6
20
Skew:                    -0.625    Prob(JB):               1.22e-
05
Kurtosis:                4.494    Cond. No.                6.
30
=====
==

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [43]:

```
x_train_lm = x_train[['carlength', 'carwidth']]
```

In [44]:

```
import statsmodels.api as sm
x_train_lm = sm.add_constant(x_train_lm)

lr = sm.OLS(y_train, x_train_lm).fit()

lr.params
```

Out[44]:

```
const      0.558465
carlength  -0.112418
carwidth    0.041778
dtype: float64
```

In [45]:

```
print(lr.summary())
```

```

=====
                        OLS Regression Results
=====
==
Dep. Variable:          stroke    R-squared:                0.0
56
Model:                  OLS      Adj. R-squared:             0.0
43
Method:                 Least Squares    F-statistic:          4.1
62
Date:                   Sun, 26 Apr 2020    Prob (F-statistic):    0.01
75
Time:                   09:11:30    Log-Likelihood:       -31.9
59
No. Observations:       143    AIC:                  69.
92
Df Residuals:           140    BIC:                  78.
81
Df Model:                2
Covariance Type:        nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.975
const	0.5585	1.330	0.420	0.675	-2.070	3.1
carlength	-0.1124	0.235	-0.478	0.633	-0.577	0.3
carwidth	0.0418	0.022	1.920	0.057	-0.001	0.0

```

-----
--
Omnibus:                15.116    Durbin-Watson:          2.0
21
Prob(Omnibus):          0.001    Jarque-Bera (JB):       20.5
46
Skew:                   -0.604    Prob(JB):               3.46e-
05
Kurtosis:               4.411    Cond. No.               3.46e+
03
=====
==

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.46e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [46]:

```
x_train_lm = x_train[['carlength', 'carwidth', 'carheight']]
```

In [47]:

```
print(lr.summary())
```

```

                                OLS Regression Results
=====
==
Dep. Variable:                stroke    R-squared:                0.0
56
Model:                        OLS      Adj. R-squared:          0.0
43
Method:                      Least Squares    F-statistic:            4.1
62
Date:                        Sun, 26 Apr 2020    Prob (F-statistic):      0.01
75
Time:                        09:11:32    Log-Likelihood:         -31.9
59
No. Observations:            143    AIC:                    69.
92
Df Residuals:                140    BIC:                    78.
81
Df Model:                    2
Covariance Type:            nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.97
const	0.5585	1.330	0.420	0.675	-2.070	3.1
carlength	-0.1124	0.235	-0.478	0.633	-0.577	0.3
carwidth	0.0418	0.022	1.920	0.057	-0.001	0.0

```

=====
==
Omnibus:                    15.116    Durbin-Watson:           2.0
21
Prob(Omnibus):              0.001    Jarque-Bera (JB):        20.5
46
Skew:                      -0.604    Prob(JB):                3.46e-
05
Kurtosis:                   4.411    Cond. No.                3.46e+
03
=====
==

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.46e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [48]:

```
cars.columns
```

Out[48]:

```
Index(['CarName', 'fueltype', 'aspiration', 'doornumber', 'carbody',  
      'drivewheel', 'enginelocation', 'wheelbase', 'carlength', 'carwidth',  
      'carheight', 'curbweight', 'enginetype', 'cylindernumber', 'enginesiz  
e',  
      'fuelsystem', 'boreratio', 'stroke', 'compressionratio', 'horsepowe  
r',  
      'peakrpm', 'citympg', 'highwaympg', 'price'],  
      dtype='object')
```

In [49]:

```
x_train_lm = x_train[['horsepower', 'boreratio', 'wheelbase']]
```

In [50]:

```
import statsmodels.api as sm  
x_train_lm = sm.add_constant(x_train_lm)  
lr = sm.OLS(y_train, x_train_lm).fit()  
lr.params
```

Out[50]:

```
const          3.002135  
horsepower     0.357183  
boreratio     -0.506372  
wheelbase      0.018691  
dtype: float64
```


In [51]:

```
print(lr.summary())
```

OLS Regression Results

```
=====
==
Dep. Variable:          stroke    R-squared:                0.1
65
Model:                  OLS      Adj. R-squared:            0.1
47
Method:                 Least Squares    F-statistic:          9.1
55
Date:                   Sun, 26 Apr 2020    Prob (F-statistic):    1.44e-
05
Time:                   09:11:41    Log-Likelihood:        -23.1
98
No. Observations:       143    AIC:                    54.
40
Df Residuals:           139    BIC:                    66.
25
Df Model:                3
Covariance Type:        nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.97
const	3.0021	0.467	6.426	0.000	2.078	3.9
horsepower	0.3572	0.174	2.054	0.042	0.013	0.7
boreratio	-0.5064	0.116	-4.354	0.000	-0.736	-0.2
wheelbase	0.0187	0.005	3.966	0.000	0.009	0.0

```
=====
==
Omnibus:                14.699    Durbin-Watson:          2.2
61
Prob(Omnibus):           0.001    Jarque-Bera (JB):        45.6
65
Skew:                    -0.171    Prob(JB):                1.21e-
10
Kurtosis:                5.747    Cond. No.                1.95e+
03
=====
==
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.95e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
#CHECKING VIF  $VIF_i = 1/(1 - R_i^2)$ 
```

In [52]:

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

In [53]:

```
def build_model(X,y):
    X = sm.add_constant(X) #Adding the constant
    lm = sm.OLS(y,X).fit() # fitting the model
    print(lm.summary()) # model summary
    return X

def checkVIF(X):
    vif = pd.DataFrame()
    vif['Features'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return(vif)
```

In [54]:

```
x_train_new = build_model(x_train_lm,y_train)
```

OLS Regression Results

```
=====
==
Dep. Variable:          stroke    R-squared:                0.1
65
Model:                  OLS      Adj. R-squared:            0.1
47
Method:                 Least Squares    F-statistic:          9.1
55
Date:                  Sun, 26 Apr 2020    Prob (F-statistic):    1.44e-
05
Time:                  09:11:48    Log-Likelihood:        -23.1
98
No. Observations:      143    AIC:                    54.
40
Df Residuals:          139    BIC:                    66.
25
Df Model:              3
Covariance Type:       nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.97
const	3.0021	0.467	6.426	0.000	2.078	3.9
horsepower	0.3572	0.174	2.054	0.042	0.013	0.7
boreratio	-0.5064	0.116	-4.354	0.000	-0.736	-0.2
wheelbase	0.0187	0.005	3.966	0.000	0.009	0.0

```
-----
--
Omnibus:              14.699    Durbin-Watson:          2.2
61
Prob(Omnibus):        0.001    Jarque-Bera (JB):        45.6
65
Skew:                 -0.171    Prob(JB):                1.21e-
10
Kurtosis:             5.747    Cond. No.                1.95e+
03
=====
==
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.95e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [55]:

```
x_train_new = build_model(x_train_new,y_train)
```

OLS Regression Results

```
=====
==
Dep. Variable:          stroke    R-squared:                0.1
65
Model:                  OLS      Adj. R-squared:            0.1
47
Method:                 Least Squares    F-statistic:           9.1
55
Date:                  Sun, 26 Apr 2020    Prob (F-statistic):     1.44e-
05
Time:                  09:11:49    Log-Likelihood:         -23.1
98
No. Observations:      143    AIC:                    54.
40
Df Residuals:          139    BIC:                    66.
25
Df Model:               3
Covariance Type:       nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.97
5]						
--						
const	3.0021	0.467	6.426	0.000	2.078	3.9
26						
horsepower	0.3572	0.174	2.054	0.042	0.013	0.7
01						
boreratio	-0.5064	0.116	-4.354	0.000	-0.736	-0.2
76						
wheelbase	0.0187	0.005	3.966	0.000	0.009	0.0
28						
=====						
==						
Omnibus:	14.699		Durbin-Watson:	2.2		
61						
Prob(Omnibus):	0.001		Jarque-Bera (JB):	45.6		
65						
Skew:	-0.171		Prob(JB):	1.21e-		
10						
Kurtosis:	5.747		Cond. No.	1.95e+		
03						
=====						
==						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.95e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [56]:

```
checkVIF(x_train_new)
```

Out[56]:

	Features	VIF
0	const	374.60
2	boreratio	1.57
1	horsepower	1.41
3	wheelbase	1.35

In [57]:

```
x_train_new = x_train_new.drop(["wheelbase"], axis = 1)
```

In [58]:

```
xtrain_new = build_model(x_train_new,y_train)
```

OLS Regression Results

```
=====
==
Dep. Variable:          stroke    R-squared:                0.0
70
Model:                  OLS      Adj. R-squared:            0.0
57
Method:                 Least Squares    F-statistic:           5.3
09
Date:                  Sun, 26 Apr 2020    Prob (F-statistic):     0.005
99
Time:                  09:12:00    Log-Likelihood:        -30.8
62
No. Observations:      143    AIC:                    67.
72
Df Residuals:          140    BIC:                    76.
61
Df Model:              2
Covariance Type:       nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.97
const	4.2675	0.359	11.895	0.000	3.558	4.9
horsepower	0.4794	0.180	2.665	0.009	0.124	0.8
boreratio	-0.3406	0.114	-2.985	0.003	-0.566	-0.1

```
-----
--
Omnibus:              5.235    Durbin-Watson:          2.1
61
Prob(Omnibus):        0.073    Jarque-Bera (JB):        7.5
92
Skew:                 0.049    Prob(JB):                0.02
25
Kurtosis:             4.124    Cond. No.                 5
2.8
=====
==
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [59]:

```
checkVIF(x_train_new)
```

Out[59]:

	Features	VIF
0	const	199.87
1	horsepower	1.37
2	boreratio	1.37

Residual analysis of a model

In [60]:


```
lm = sm.OLS(y_train,x_train_new).fit()  
y_train_price = lm.predict(x_train_new)
```

In [61]:

```
fig = plt.figure()  
sns.distplot((y_train - y_train_price), bins = 20)  
fig.suptitle('Error Terms', fontsize = 20)  
plt.xlabel('Errors', fontsize = 18)
```

Out[61]:

```
Text(0.5, 0, 'Errors')
```



MODEL EVALUATION

In [62]:

```
num_vars = ['wheelbase', 'curbweight', 'enginesize', 'carlength', 'carwidth', 'price']  
x_train[num_vars] = scaler.transform(x_train[num_vars])
```

In [63]:

```
y_test = x_train.pop('price')  
x_test = x_train
```

In [64]:

```
x_train_new = x_train_new.drop('const',axis=1)  
x_test_new = x_test[x_train_new.columns]  
x_test_new = sm.add_constant(x_test_new)
```

In [65]:

```
y_pred = lm.predict(x_test_new)
```

In [66]:

```
from sklearn.metrics import r2_score  
r2_score(y_test, y_pred)
```

Out[66]:


-326434617006.14984

In [67]:

```
fig = plt.figure()  
plt.scatter(y_test,y_pred)  
fig.suptitle('y_test vs y_pred', fontsize=20)  
plt.xlabel('y_test', fontsize=18)  
plt.ylabel('y_pred', fontsize=16)
```

Out[67]:

Text(0, 0.5, 'y_pred')



In [68]:

```
print(lm.summary())
```

OLS Regression Results						
=====						
==						
Dep. Variable:	stroke	R-squared:	0.070			
Model:	OLS	Adj. R-squared:	0.057			
Method:	Least Squares	F-statistic:	5.309			
Date:	Sun, 26 Apr 2020	Prob (F-statistic):	0.00599			
Time:	09:12:45	Log-Likelihood:	-30.862			
No. Observations:	143	AIC:	67.72			
Df Residuals:	140	BIC:	76.61			
Df Model:	2					
Covariance Type:	nonrobust					
=====						
==						
	coef	std err	t	P> t	[0.025	0.975]

--						
const	4.2675	0.359	11.895	0.000	3.558	4.977
horsepower	0.4794	0.180	2.665	0.009	0.124	0.835
boreratio	-0.3406	0.114	-2.985	0.003	-0.566	-0.115
=====						
==						
Omnibus:	5.235	Durbin-Watson:	2.161			
Prob(Omnibus):	0.073	Jarque-Bera (JB):	7.592			
Skew:	0.049	Prob(JB):	0.025			
Kurtosis:	4.124	Cond. No.	52.8			
=====						
==						

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In []:

In []: